

• • LIVE WEEKEND COHORT

Build with AI — Weekend Cohort

Go from AI consumer to AI builder in one weekend. Ship three working AI-powered applications using real APIs, RAG, and agent patterns.

PRICE	DURATION	LEVEL	FORMAT	SCHEDULE
₹2,499	2 Days · 14 hrs	Beginner–Intermediate	Live	Bi-weekly weekends



Prasant Mishra

Founder & Instructor, Qurious Academy

What is Build with AI — Weekend Cohort?

Most people who are excited about AI are still just using it — copying prompts, trying different tools, watching demos. This weekend flips that. By Sunday evening you'll have moved from consumer to builder: you'll have written real code that calls AI APIs, implemented a RAG (Retrieval-Augmented Generation) system, built an AI writing assistant with memory, and deployed your own AI micro-product.

The cohort is structured around projects, not lectures. You learn a concept, immediately build with it, then move to the next. Day 1 takes you from understanding how LLMs work through your first two projects. Day 2 covers advanced patterns — function calling, agents, embeddings — and ends with you building and deploying your own AI project concept.

This is a hands-on coding cohort. You'll need basic Python (variables, functions, loops) but no AI knowledge at all. The emphasis throughout is on patterns: the code patterns that appear in almost every AI application, so you can adapt them to any project you want to build next.

IS THIS RIGHT FOR YOU?

Who Should Attend

✓ THIS IS FOR YOU IF...

- ✓ You know basic Python and want to start building with AI — not just understanding it theoretically
- ✓ You're a developer who hasn't worked with AI yet and wants to get up to speed quickly with practical patterns
- ✓ You have a product or business idea that could use AI and want to validate/prototype it over a weekend
- ✓ You want to add AI skills to your professional profile in the shortest possible time

✗ THIS IS NOT FOR YOU IF...

- ✗ You don't know Python at all — take the Python Weekend Intensive first, then come back to this
- ✗ You're looking for deep ML theory (training models, backpropagation, PyTorch) — see the AI & ML Engineering 12-Week program
- ✗ You want to build production-scale AI systems — this cohort focuses on getting to working prototypes, not production engineering

Prerequisites: Basic Python — you should be comfortable with variables, functions, and loops. No AI, ML, or API experience needed. We'll cover everything from there.

Two Days, Four Blocks, Three Projects

DAY 1 — MORNING

AI Foundations & APIs

- ▶ **How LLMs work (intuition)** — a fast, visual explanation of transformers, tokens, and context windows; just enough to make sense of what you're building with
- ▶ **OpenAI & Anthropic API setup** — accounts, API keys, rate limits, pricing; everything you need to start calling LLMs in code
- ▶ **Prompt engineering** — zero-shot, few-shot, and chain-of-thought prompting with live examples of what works and what doesn't; system vs. user prompts
- ▶ **Structured output & JSON mode** — getting predictable, machine-parseable output from LLMs; the foundation of every serious AI application

DAY 1 — AFTERNOON

Projects 1 & 2

- ▶ **Project 1: AI Q&A over your own documents (RAG basics)** — implement a simple Retrieval-Augmented Generation system; upload documents, embed them, and answer questions with citations
- ▶ **Project 2: AI writing assistant with memory** — build a conversational AI that remembers context across turns; understanding conversation history and memory patterns
- ▶ **Token limits, costs & best practices** — understanding context windows, how to manage token usage, and keeping API costs under control in real applications
- ▶ **Error handling & fallbacks** — building AI applications that degrade gracefully; retry logic, fallback models, and handling rate limits

DAY 2 — MORNING

Advanced Patterns

- ▶ **Function calling & tool use** — giving LLMs the ability to call your code; the pattern behind AI assistants that can take actions in the real world
- ▶ **Agents — multi-step reasoning** — building AI systems that can plan, execute steps, and respond to intermediate results; the ReAct pattern explained
- ▶ **Embeddings & semantic search** — vector representations of text; building search that understands meaning, not just keywords
- ▶ **Evaluating AI outputs** — how to measure whether your AI application is actually working; LLM-as-judge, human eval, and automated metrics

DAY 2 — AFTERNOON

Project 3 & Deployment

- ▶ **Project 3: your own AI micro-SaaS concept** — you design and build a small AI application around your own use case; the instructor provides live technical support
- ▶ **Deploying AI apps (Vercel + API keys)** — shipping your application to a live URL; environment variables, secret management, and the deployment checklist
- ▶ **Cost optimisation** — caching, prompt compression, model selection, and batching strategies to keep running costs low

- ▶ **What to build next — personal roadmap** — based on your Project 3, a concrete next-steps plan for taking your idea further

LEARNING OUTCOMES

What You'll Walk Away With

01 Build 3 working AI-powered applications
— a RAG system, a conversational assistant, and your own original AI project
— all deployed and shareable

02 Understand prompt engineering deeply
— not surface-level tips but a systematic understanding of how to shape LLM behaviour reliably

03 Implement RAG and agent patterns
— the two most important architectural patterns in modern AI applications, implemented from scratch

04 Deploy an AI product end-to-end
— from local development to a live URL with proper API key management and cost controls

WHAT'S INCLUDED

Everything You Get

✓ **Both day recordings** — full HD recordings of both sessions, shared within 24 hours

✓ **All project code** — complete, annotated code for all three projects as starter templates

✓ **Prompt template library** — a collection of reusable, battle-tested prompt templates for common AI application patterns

✓ **2 weeks doubt support** — post-cohort WhatsApp access to ask implementation questions

✓ **Learner's Certificate** — digital certificate of completion from Qurious Academy

LOGISTICS

Schedule & Pricing

SCHEDULE

Every other weekend (bi-weekly)

TIME

9:00 AM – 4:00 PM IST

DURATION

2 days · ~14 hours total

PRICE

₹2,499

❏ **Will I need to pay for API access separately?**

Yes — you'll need an OpenAI or Anthropic account. For the weekend's projects, API costs are typically under ₹200–400. We'll show you how to set spending limits to control costs.

❏ **What Python level do I need?**

You should be comfortable with variables, functions, loops, and importing libraries. You don't need to know classes, decorators, or anything advanced.

❏ **Can I attend if I use JavaScript instead of Python?**

The cohort teaches Python. JavaScript developers can follow the concepts, but the code examples will all be in Python. We recommend picking up basic Python before attending.

About Prasant Mishra



Prasant Mishra

Founder & Instructor — Qurious Academy

GenAI Solutions Architect

15+ Yrs Teaching

2500+ Students

60+ Projects Delivered

Prasant designs and ships AI-powered systems as a GenAI Solutions Architect. He has built RAG pipelines, AI agents, and LLM-powered products for real clients — which means the patterns he teaches are battle-tested in production, not theoretical constructs from an academic course.

He brings that same practical, project-first mindset to this cohort: less theory, more code, real output by the end of the weekend. His goal isn't for you to understand AI — it's for you to be able to build with it.

QuriousAcademy

Focused, live learning for curious minds. No fluff — just the concepts that matter.

Build with AI — Weekend Cohort

Live Cohort · Bi-weekly weekends · 9 AM–4 PM IST

₹2,499

Enroll at quriousacademy.com

or visit quriousacademy.com/courses

Questions? Reach us at hello@quriousacademy.com

© 2025 Qurious Academy. All rights reserved.